



elasticsearch

By : - MAKOUNDOU Paule-Theodora
DE TROGOFF Guilhem
SHARMA Karan
LEDROLE Pierre

What is ElasticSearch h ?

-
- Elasticsearch is a BASE (Basically Available, Soft State, Eventually Consistent) system.
 - Prioritizes availability and scalability over strict consistency.
 - Stores and indexes JSON documents for fast search and retrieval





Examples of Data Stored in Elasticsearch

- **Logs and Event Data:** System and application logs.

- **Full-Text Documents:** Articles, product descriptions.

- **Structured Data:** User information, product details.

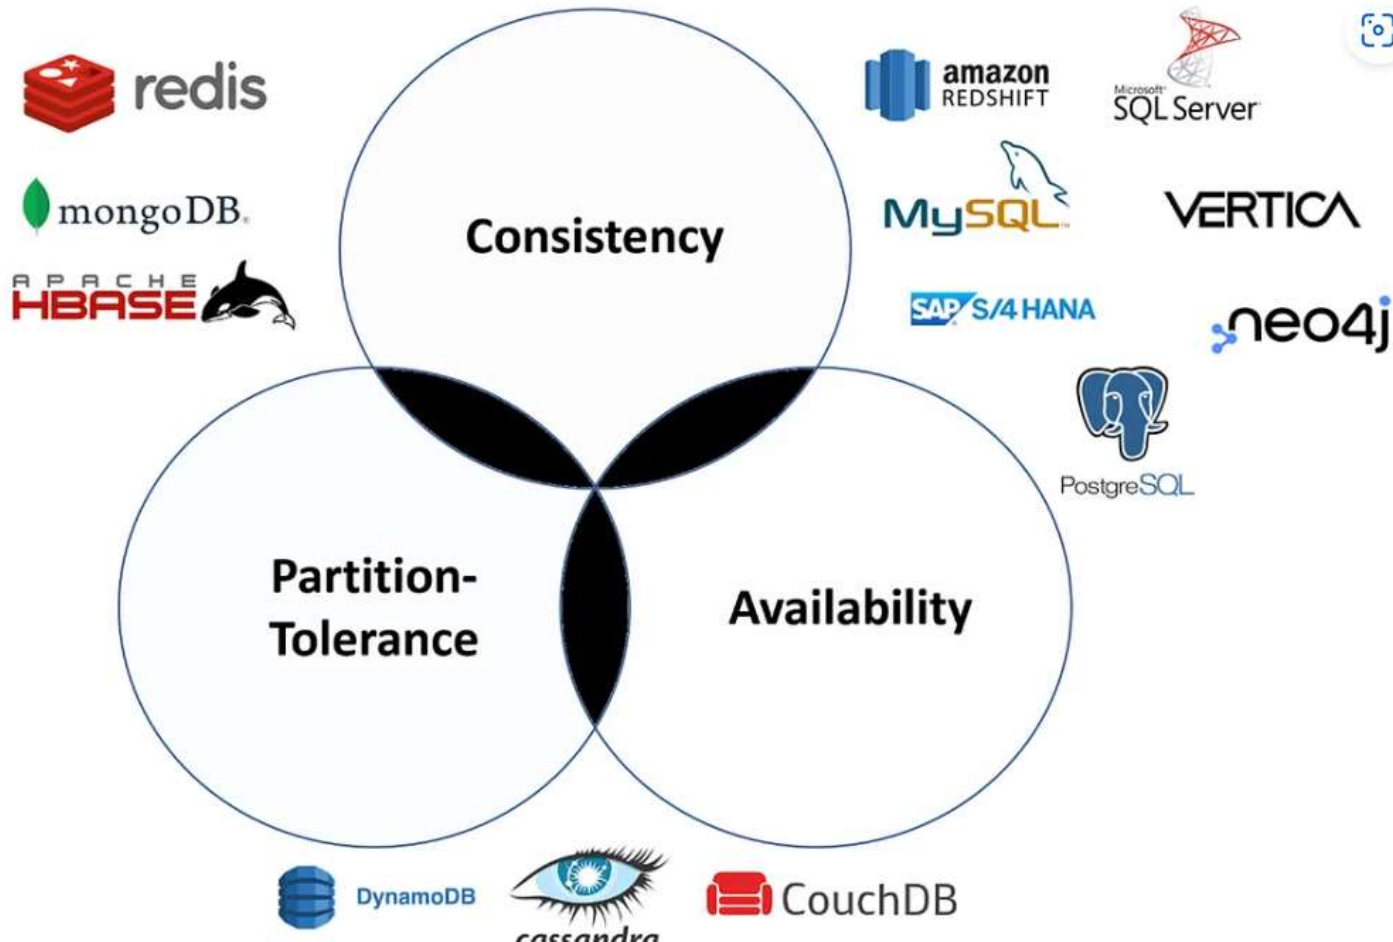
- **Geospatial Data:** Geographic coordinates, maps.

- **Time-Series Data:** Metrics, financial data.

- **Social Media Data:** Posts, user interactions.

- **Machine Learning Data:** Training sets, predictions.

CAP theorem positioning



- **Consistency:**
Ensures data across nodes is synchronized.
- **Partition Tolerance:**
Continues operating despite network issues.
- **Availability:**
Sacrifices some availability during network partitions to maintain consistency.

Examples of Uses

- **Uses:** Used by companies like Netflix and GitHub. Powers search engines for platforms like Wikipedia and Shopify.
- **Issues:** High resource usage (RAM, CPU), complex cluster management, risk of data loss from poor configuration.
- **Cost:** Free for the open-source version; paid options offer advanced features. Managed Elastic Cloud services simplify setup but increase costs



How Popular is this Tool?

- Elasticsearch is widely popular, especially in IT, e-commerce, and media.
- Major companies like Netflix, Uber, and Shopify rely on it
- Strong community support, frequent updates, and adaptability for real-time, large-scale search needs.

The Splunk logo, featuring the word "splunk" in a black, lowercase, sans-serif font, followed by a green greater-than sign (>).The Amazon logo, featuring the word "amazon" in a black, lowercase, sans-serif font, with a curved orange arrow underneath it.The Apache Solr logo, featuring the word "Apache" in a black, sans-serif font above the word "Solr" in a larger, bold, black, sans-serif font. To the right of the text is a circular icon composed of many thin, radiating lines in shades of orange and yellow.

Main Competitors

Competitors: Apache Solr (open-source), Splunk (commercial), Amazon CloudSearch (managed), and Lucene for custom search.

Timeline: Released in 2010, evolved with ELK integration, cloud offerings, and machine learning features.

Future and Market Changes



ElasticSearch is
expected to expand:



in machine learning



in cloud services.


Benefits and limitations

- Benefits: Highly scalable, real-time search and analytics, flexible full-text search capabilities, open-source with strong integration options.
- Limitations: No ACID compliance (BASE model), complex cluster management, high resource demands increase operational costs.




Indexing in Elasticsearch and its code implementation

- Indexing in Elasticsearch is essentially the process of storing and organizing data so that it can be efficiently searched and retrieved. When we index a document, we add it to an **index** (like a table in SQL) where it becomes searchable.
- **Index:** A collection of documents (like a table in SQL).
- **Document:** The basic unit of information (like a row in SQL) containing various fields (key-value pairs).
- **Field:** Each key-value pair in a document (like columns in SQL).
- **Shards:** Indexes are split into smaller units called shards, which allow Elasticsearch to distribute data across multiple nodes for scalability and performance.



The Indexing Process:

- 1. Parsing:** Elasticsearch breaks down a document into its individual fields during indexing.
- 2. Mapping:** Elasticsearch automatically maps field types (or you can define your own mappings), telling Elasticsearch how to store and search the data.
- 3. Inverted Index:** For text fields, Elasticsearch creates an inverted index that allows for fast full-text search by



Elasticsearch Python

```
1 from elasticsearch import Elasticsearch
2
3 # Initialize the Elasticsearch client
4 es = Elasticsearch("http://localhost:9200")
5
6 # Index a document
7 doc = {
8     'name': 'Karan',
9     'age': 24,
10    'occupation': 'Developer'
11 }
12 es.index(index='users', id=1, document=doc)
```

Code Example (Python) :

Indexing a
document into
Elasticsearch:

Similarities with SQL and Code Comparison

Similarities with SQL:

- **Indexes vs. Tables:** In Elasticsearch, **indexes** are similar to **tables** in SQL. Both are used to store collections of documents (rows).
- **Documents vs. Rows:** In Elasticsearch, a **document** is like a **row** in SQL.
- **Fields vs. Columns:** Elasticsearch's **fields** correspond to **columns** in SQL.
- **Queries:** Both Elasticsearch and SQL allow filtering and searching based on conditions.
- **Aggregations:** Elasticsearch **aggregations** are similar to **SQL aggregate functions** (e.g., COUNT, SUM, AVG).

Key Differences:



Full-Text Search:

Elasticsearch excels at full-text search, while SQL is not optimized for it.



Schema: Elasticsearch is schema-less, whereas SQL databases require a predefined schema.



Joins: SQL supports complex joins between tables, while Elasticsearch is document-based and does not support traditional joins.

Code Comparison (Examples) :

Insert a row:

Elasticsearch Python

```
1 from elasticsearch import Elasticsearch
2
3 # Initialize the Elasticsearch client
4 es = Elasticsearch("http://localhost:9200")
5
6 # Index a document
7 doc = {
8     'name': 'Karan',
9     'age': 24,
10    'occupation': 'Developer'
11 }
12 es.index(index='users', id=1, document=doc)
```

SQL

```
1 INSERT INTO users (id, name, age, occupation)
2 VALUES (1, 'Karan', 24, 'Developer');
```

Query with conditions:

Elasticsearch Python

```
1 es.search(index='users', query={
2     "bool": {
3         "must": [
4             {"range": {"age": {"gt": 25}}},
5             {"match": {"occupation": "Developer"}}
6         ]
7     }
8 })
```

SQL

```
1 SELECT * FROM users WHERE age > 25 AND occupation = 'Developer';
```


Aggregate (COUNT) and group by:

Elasticsearch Python

```
1 es.search(index='users', size=0, aggs={  
2     "group_by_occupation": {  
3         "terms": {"field": "occupation.keyword"}  
4     }  
5 })
```


SQL

```
1 SELECT occupation, COUNT(*) FROM users GROUP BY occupation;
```

Scalability:

Elasticsearch is designed to be highly scalable. As your data grows, Elasticsearch can scale horizontally by adding more nodes to your cluster. Key features like **sharding** and **replication** make it resilient and able to handle large volumes of data efficiently.

- **Sharding:** Each index is divided into shards, which are distributed across multiple nodes.
- **Replication:** Shards are replicated to provide high availability. If a node goes down, Elasticsearch will use replica shards to ensure data availability.



Flexibility

One of Elasticsearch's greatest strengths is its flexibility. It handles structured, semi-structured, and unstructured data seamlessly. You can store:

- **Text** and perform full-text search.
- **Numbers, dates, and booleans** for structured queries.
- **Arrays** and **nested objects** for more complex data models.
- **Geospatial data** for location-based searches.
- Elasticsearch is also **schema-less** by default, meaning you don't need to define a rigid structure upfront. You can index and search a wide variety of data types without setting up predefined mappings (though you can define mappings if needed).

Usability



RESTful API: You can interact with Elasticsearch using simple HTTP requests (e.g., GET, POST), making it easy to use and integrate with different tools.



Kibana: Kibana, part of the Elastic Stack, allows you to visualize and explore your data using a graphical interface. This makes Elasticsearch accessible to technical and non-technical users alike.



Real-Time Data Ingestion: Elasticsearch indexes data in near real-time, meaning new data becomes searchable almost immediately after it's ingested.

